

Multi-Object Tracking Project

Outline

- Introduction
- Project Overview & Objectives
- Design Process & Challenges
- Implementation Details & Challenges
- Demo
- Conclusion

Introduction




Introduction




Introduction

Project timeframe!



<https://mcscert.mcmaster.ca/>

Overview



Overview

Could we perform object **tracking** and **classification** with a
2D LiDAR and an RGB camera?

Overview

Could we perform object **tracking** and **classification** with a
2D LiDAR and an RGB camera?



Track 1
Track 2
...
Track N

Track {
id
position
velocity
class
...
}

Objective

To **design** and **implement** a multi-object tracking
and classification pipeline in the RC car

Objective

To **design and implement** a multi-object tracking
and classification pipeline in the RC car

To **learn as much as I can** about perception, and
demonstrate my learning

Objective

To **design and implement** a multi-object tracking
and classification pipeline in the RC car

To **learn as much as I can** about perception, and
demonstrate my learning

4-months

Objective

To **design and implement** a multi-object tracking
and classification pipeline in the RC car

To **learn as much as I can** about perception, and
demonstrate my learning

part-time

4-months

Objective

To **design and implement** a multi-object tracking
and classification pipeline in the RC car

To **learn as much as I can** about perception, and
demonstrate my learning

part-time

4-months

little robotics
experience

Objective

To **design and implement** a multi-object tracking
and classification pipeline in the RC car

To **learn as much as I can** about perception, and
demonstrate my learning

| | | | |
|----------|-----------|----------------------------|---|
| 4-months | part-time | little robotics experience | efficient implementation get SOMETHING working |
|----------|-----------|----------------------------|---|

Part 1: Design

Sensor Input

2D-LiDAR Scans

$$z_i = \{d_i, \theta_i\}$$

$$Z = \{z_i\}_{i=1}^L$$

**Monocular RGB
Images**



Sensor Input

2D-LiDAR Scans

$$z_i = \{d_i, \theta_i\}$$

$$Z = \{z_i\}_{i=1}^L$$


**Monocular RGB
Images**



Output

Labelled Tracks

Track {
id
position
velocity
class
...}



Design Challenge 1: Knowledge gap

2D-LiDAR Scan

$$z_i = \{d_i, \theta_i\}$$

$$Z = \{z_i\}_i^L$$

Monocular RGB Images





LiDAR Pipeline


Output


What is the pinhole camera model?
What does it mean to calibrate a camera?
Homogenous coordinates?
Transformation matrices?
...
+ more stuff I had to learn as I was designing

Modelled Tracks
ack {
id
position
velocity
class
...}










Design Challenge 2: Sparse LiDAR points further away

2D-LiDAR Scan

$$z_i = \{d_i, \theta_i\}$$

$$Z = \{z_i\}_i^L$$

$$r_i = r_0 + d_i * constant$$




Monocular RGB Images



Output

Detected Tracks

```
track {  
    id  
    position  
    velocity  
    class  
    ...}
```



Sensor Input

LiDAR Pipeline

Output

2D-LiDAR Scan

$$z_i = \{d_i, \theta_i\}$$

$$Z = \{z_i\}_{i=1}^L$$

Monocular RGB Images



L-Shape Fitting [1]

Algorithm 2 Search-Based Rectangle Fitting


Input: range data points $X \in R^{n \times 2}$

Output: rectangle edges $\{a_i x + b_i y = c_i | i = 1, 2, 3, 4\}$

```

1:  $Q \leftarrow \emptyset$ 
2: for  $\theta = 0$  to  $\pi/2 - \delta$  step  $\delta$  do
3:    $\hat{e}_1 \leftarrow (\cos \theta, \sin \theta)$      $\triangleright$  rectangle edge direction vector
4:    $\hat{e}_2 \leftarrow (-\sin \theta, \cos \theta)$ 
5:    $C_1 \leftarrow X \cdot \hat{e}_1^T$             $\triangleright$  projection on to the edge
6:    $C_2 \leftarrow X \cdot \hat{e}_2^T$ 
7:    $q \leftarrow \text{CalculatecriterionX}(C_1, C_2)$ 
8:   insert  $q$  into  $Q$  with key  $(\theta)$ 
9: end for
10: select key  $(\theta^*)$  from  $Q$  with maximum value
11:  $C_1^* \leftarrow X \cdot (\cos \theta^*, \sin \theta^*)^T, C_2^* \leftarrow X \cdot (-\sin \theta^*, \cos \theta^*)^T$ 
12:  $a_1 \leftarrow \cos \theta^*, b_1 \leftarrow \sin \theta^*, c_1 \leftarrow \min\{C_1^*\}$ 
13:  $a_2 \leftarrow -\sin \theta^*, b_2 \leftarrow \cos \theta^*, c_2 \leftarrow \min\{C_2^*\}$ 
14:  $a_3 \leftarrow \cos \theta^*, b_3 \leftarrow \sin \theta^*, c_3 \leftarrow \max\{C_1^*\}$ 
15:  $a_4 \leftarrow -\sin \theta^*, b_4 \leftarrow \cos \theta^*, c_4 \leftarrow \max\{C_2^*\}$ 


```





labelled Tracks
ack {
id
position
velocity
class
...}

Search based algorithm to fit a rectangle to a cluster of points


[1] X. Zhang, W. Xu, C. Dong and J. M. Dolan, "Efficient L-shape fitting for vehicle detection using laser scanners," 2017 IEEE Intelligent Vehicles Symposium (IV), Los Angeles, CA, USA, 2017, pp. 54-59, doi: 10.1109/IVS.2017.7995698. keywords: {Feature extraction;Shape;Laser radar;Three-dimensional displays;Clustering algorithms;Vehicle detection;Automobiles},








Design Challenge 3: Large jumps in rectangle




Design Challenge 3: Large jumps in rectangle



Solution 1: Use the largest length & width between the rectangles, recalculate center
e.g. `curr_frame.length() = max(curr_frame.length(), prev_frame.length())`


Design Challenge 3: Large jumps in rectangle



Solution 1: Use the largest length & width between the rectangles, recalculate center
e.g. `curr_frame.length() = max(curr_frame.length(), prev_frame.length())`


But what if two objects get fused together??

Design Challenge 3: Large jumps in rectangle



Solution 2: Keep a circular queue of last n associated detections, take the median


Design Challenge 3: Large jumps in rectangle





Solution 3: Use a Kalman Filter to estimate the length and width

Design Challenge 3.1: Identity switching

Did not consider








Design Challenge 4: Asynchronous LiDAR and Camera streams

Camera frequency: 15 – 30 fps

LiDAR frequency: 40 Hz



Idea 1: When a track gets assigned a label, it keeps the label until new camera data comes in


Only apply a label to a track if it has been classified consistently

Association back in time: keep circular queue of previous tracks & timestamps, associate to those within 10ms? `queue.length() = # of timesteps` camera lags behind

Design Challenge 4: Asynchronous LiDAR and Camera streams

Camera frequency: 15 – 30 fps

LiDAR frequency: 40 Hz





Idea 2: Somehow interpolate the tracks

Part 2: Implementation

Implementation



GAZEBO




Implementation Challenge 1:

Problem: Learning curve! First time using ROS.

Solution: Looking at documentation, examples, trying and failing!!


Implementation Challenge 2:

Problem: In the past, students did not have a streamlined way to develop software.




Implementation Challenge 2:

Problem: In the past, students did not have a streamlined way to develop software.



Idea: Using Docker for more streamlined development

Downsides: GPIO trouble, Docker overhead




Implementation Challenge 3:

Problem: How do I know it's working? How do we define "working"?

Implementation Challenge 3:

Problem: How do I know it's working? How do we define "working"?


In this project: Modular development and testing w/ ROS bag & simple Gazebo environment



Implementation Challenge 3:

Problem: How do I know it's working? How do we define "working"?


In this project: Modular development and testing w/ ROS bag & simple Gazebo environment



Evaluate? Visually. Not ideal, but enough for this project scope.



Ideally: Evaluate GOSPA against ground truth + benchmarking

Demo



Implementation Challenge 4:

Problem: Unconsidered behaviour of LShapeFitting algorithm



Output from LShapeFitting:





$$W = 0.3\text{m}$$

$$L = 0.7\text{m}$$

$$\theta = 0 \text{ deg}$$

Implementation Challenge 4:

Problem: Unconsidered behaviour of LShapeFitting algorithm



Output from LShapeFitting:



$W = 0.3m$
 $L = 0.7m$
 $\theta = 0 \text{ deg}$



OR

$W = 0.7m$
 $L = 0.3m$
 $\theta = 90 \text{ deg}$

??

Implementation Challenge 4:

Problem: Unconsidered behaviour of LShapeFitting algorithm



Output from LShapeFitting:



$W = 0.3m$
 $L = 0.7m$
 $\theta = 0 \text{ deg}$

OR

$W = 0.7m$
 $L = 0.3m$
 $\theta = 90 \text{ deg}$

UNSOLVED!!

Implementation Challenge 5:

Problem: Finding the extrinsic parameters between the 2D-LiDAR and the monocular camera. The lack of a third dimension makes it hard to find point-pixel correspondences.

Solution: Had some ideas, but ultimately never tried as the camera pipeline was never implemented.

Thank you for listening!

References

- [1] X. Zhang, W. Xu, C. Dong and J. M. Dolan, "Efficient L-shape fitting for vehicle detection using laser scanners," 2017 IEEE Intelligent Vehicles Symposium (IV), Los Angeles, CA, USA, 2017, pp. 54-59, doi: 10.1109/IVS.2017.7995698. keywords: {Feature extraction;Shape;Laser radar;Three-dimensional displays;Clustering algorithms;Vehicle detection;Automobiles},